

Predicting an NBA Player's Position

Project Goal

The NBA is considered the most "positionless" major sport. The game has evolved to the point where 6'8" players could play any position, from guard to center. Can we use machine learning to predict an NBA player's position based on their in-game stats?

Data

NBA Player's Stats since 1950

Libraries Used

- Pandas
- Numpy
- Sklearn
- Seaborn
- Geopandas

Machine Learning Algorithms Used

- Decision Trees
- Random Forest
- K-Nearest Neighbors
- Support Vector Machine

1. Reading/Cleaning Data

```
In [120... import pandas as pd
pd.set_option('display.max_columns', 500)

df = pd.read_csv("Seasons_Stats.csv")

df.head()
```

Out[120...

	Unnamed: 0	Year	Player	Pos	Age	Tm	G	GS	MP	PER	TS%	3PAr	FT%	ORB%	I
0	0	1950.0	Curly Armstrong	G-F	31.0	FTW	63.0	NaN	NaN	NaN	0.368	NaN	0.467	NaN	
1	1	1950.0	Cliff Barker	SG	29.0	INO	49.0	NaN	NaN	NaN	0.435	NaN	0.387	NaN	
2	2	1950.0	Leo Barnhorst	SF	25.0	CHS	67.0	NaN	NaN	NaN	0.394	NaN	0.259	NaN	
3	3	1950.0	Ed Bartels	F	24.0	TOT	15.0	NaN	NaN	NaN	0.312	NaN	0.395	NaN	

	Unnamed: 0	Year	Player	Pos	Age	Tm	G	GS	MP	PER	TS%	3PAr	FTr	ORB%	I
4	4	1950.0	Ed Bartels	F	24.0	DNN	13.0	NaN	NaN	NaN	0.308	NaN	0.378	NaN	

```
In [121...] df.columns
```

```
Out[121...] Index(['Unnamed: 0', 'Year', 'Player', 'Pos', 'Age', 'Tm', 'G', 'GS', 'MP', 'PER', 'TS%', '3PAr', 'FTr', 'ORB%', 'DRB%', 'TRB%', 'AST%', 'STL%', 'BLK%', 'TOV%', 'USG%', 'blan1', 'OWS', 'DWS', 'WS', 'WS/48', 'blank2', 'OBPM', 'DBPM', 'BPM', 'VORP', 'FG', 'FGA', 'FG%', '3P', '3PA', '3P%', '2P', '2PA', '2P%', 'eFG%', 'FT', 'FTA', 'FT%', 'ORB', 'DRB', 'TRB', 'AST', 'STL', 'BLK', 'TOV', 'PF', 'PTS'], dtype='object')
```

Finding columns with null values

```
In [122...] df.drop(columns = ['Unnamed: 0'], inplace=True)
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 24691 entries, 0 to 24690
Data columns (total 52 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   Year        24624 non-null  float64
1   Player      24624 non-null  object
2   Pos         24624 non-null  object
3   Age         24616 non-null  float64
4   Tm          24624 non-null  object
5   G           24624 non-null  float64
6   GS          18233 non-null  float64
7   MP          24138 non-null  float64
8   PER         24101 non-null  float64
9   TS%         24538 non-null  float64
10  3PAr        18839 non-null  float64
11  FTr         24525 non-null  float64
12  ORB%        20792 non-null  float64
13  DRB%        20792 non-null  float64
14  TRB%        21571 non-null  float64
15  AST%        22555 non-null  float64
16  STL%        20792 non-null  float64
17  BLK%        20792 non-null  float64
18  TOV%        19582 non-null  float64
19  USG%        19640 non-null  float64
20  blan1       0 non-null      float64
21  OWS         24585 non-null  float64
22  DWS         24585 non-null  float64
23  WS          24585 non-null  float64
24  WS/48       24101 non-null  float64
25  blank2      0 non-null      float64
26  OBPM        20797 non-null  float64
27  DBPM        20797 non-null  float64
28  BPM         20797 non-null  float64
29  VORP        20797 non-null  float64
30  FG          24624 non-null  float64
31  FGA         24624 non-null  float64
32  FG%         24525 non-null  float64
33  3P          18927 non-null  float64
34  3PA         18927 non-null  float64
35  3P%         15416 non-null  float64
36  2P          24624 non-null  float64
```

```

37 2PA      24624 non-null float64
38 2P%      24496 non-null float64
39 eFG%     24525 non-null float64
40 FT       24624 non-null float64
41 FTA      24624 non-null float64
42 FT%      23766 non-null float64
43 ORB      20797 non-null float64
44 DRB      20797 non-null float64
45 TRB      24312 non-null float64
46 AST      24624 non-null float64
47 STL      20797 non-null float64
48 BLK      20797 non-null float64
49 TOV      19645 non-null float64
50 PF       24624 non-null float64
51 PTS      24624 non-null float64
dtypes: float64(49), object(3)
memory usage: 9.5+ MB

```

Finding duplicated rows

```
In [123...] df.duplicated().sum()
```

```
Out[123...] 66
```

```
In [124...] df['MPG'] = df['MP'] / df['G']
```

```
In [125...] df.isna().sum()
```

```

Out[125...] Year      67
Player      67
Pos         67
Age         75
Tm          67
G           67
GS          6458
MP          553
PER         590
TS%         153
3PAr        5852
FTr         166
ORB%        3899
DRB%        3899
TRB%        3120
AST%        2136
STL%        3899
BLK%        3899
TOV%        5109
USG%        5051
blan1       24691
OWS         106
DWS         106
WS          106
WS/48       590
blank2      24691
OBPM        3894
DBPM        3894
BPM         3894
VORP        3894
FG          67
FGA         67
FG%         166
3P          5764
3PA         5764

```


3P% 9275
2P 67
2PA 67
2P% 195
eFG% 166
FT 67
FTA 67
FT% 925
ORB 3894
DRB 3894
TRB 379
AST 67
STL 3894
BLK 3894
TOV 5046
PF 67
PTS 67
MPG 553
dtype: int64

```
In [126... df[df['FT%'].isnull()]
```

Out[126...

	Year	Player	Pos	Age	Tm	G	GS	MP	PER	TS%	3PAr	FTr	ORB%	DRB%
79	1950.0	Normie Glick	F	22.0	MNL	1.0	NaN	NaN	NaN	1.000	NaN	0.0	NaN	NaN
132	1950.0	Lee Knorek	C	28.0	BLB	1.0	NaN	NaN	NaN	0.000	NaN	0.0	NaN	NaN
175	1950.0	Murray Mitchell	C	26.0	AND	2.0	NaN	NaN	NaN	0.333	NaN	0.0	NaN	NaN
187	1950.0	Jim Nolan	C	22.0	PHW	5.0	NaN	NaN	NaN	0.190	NaN	0.0	NaN	NaN
312	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
...
24568	2017.0	Damjan Rudez	SF	30.0	ORL	45.0	0.0	314.0	6.3	0.466	0.727	0.0	1.7	7.1
24620	2017.0	Mike Tobey	C	22.0	CHO	2.0	0.0	25.0	-0.1	0.250	0.000	0.0	8.7	4.4
24622	2017.0	Axel Toupane	SF	24.0	TOT	4.0	0.0	47.0	6.2	0.611	0.444	0.0	0.0	2.3
24623	2017.0	Axel Toupane	SF	24.0	MIL	2.0	0.0	6.0	-9.9	0.000	1.000	0.0	0.0	0.0
24624	2017.0	Axel Toupane	SF	24.0	NOP	2.0	0.0	41.0	8.6	0.688	0.375	0.0	0.0	2.6

925 rows × 53 columns



```
In [127... df[df['FT%'].isnull()].mean()['G']
```

Out[127... 4.301864801864802

```
In [128... df[df['FT%'].isnull()].mean()['MP']
```

```
Out[128... 26.24970691676436
```

```
In [129... df[df['FT%'].isnull()].mean()['MPG']
```

```
Out[129... 5.798066690777163
```

```
In [130... all_years = df.copy()
df = df[df['Year'] >= 2010]
df = df[(df['G'] >= 20) & (df['MPG'] >= 15)]
df.shape
```

```
Out[130... (2845, 53)
```

```
In [131... df.fillna(0, inplace=True)
```

Finding the different positions in the dataset

Some positions have multiple values, such as "PF-SF". To fix this, I filtered the dataset by the specific values, and changed the positions to make it uniform.

```
In [132... df['Pos'].value_counts()
```

```
Out[132... SG      601
PG      590
PF      568
SF      566
C       481
PF-SF     7
SG-PG     7
SF-PF     6
PG-SG     6
SG-SF     4
SF-SG     3
PF-C       3
C-PF       2
SG-PF     1
Name: Pos, dtype: int64
```

```
In [133... df[df['Pos'] == 'PF-SF']
```

```
Out[133...
```

	Year	Player	Pos	Age	Tm	G	GS	MP	PER	TS%	3PAr	FTr	ORB%	DRB%
20216	2010.0	Jared Jeffries	PF-SF	28.0	TOT	70.0	37.0	1794.0	9.8	0.507	0.226	0.382	8.1	10.3
20401	2010.0	James Singleton	PF-SF	28.0	TOT	57.0	3.0	977.0	11.4	0.452	0.146	0.244	12.6	20.0
20725	2011.0	Danilo Gallinari	PF-SF	22.0	TOT	62.0	60.0	2104.0	15.7	0.597	0.458	0.611	3.2	13.4
20753	2011.0	Jeff Green	PF-SF	24.0	TOT	75.0	51.0	2427.0	12.9	0.538	0.257	0.298	3.7	13.5
22048	2013.0	Marcus Morris	PF-SF	23.0	TOT	77.0	23.0	1524.0	11.3	0.516	0.443	0.222	6.0	14.6

	Year	Player	Pos	Age	Tm	G	GS	MP	PER	TS%	3PAr	FTr	ORB%	DRB%
22625	2014.0	Luc Mbah	PF-SF	27.0	TOT	64.0	7.0	1003.0	8.3	0.503	0.088	0.347	6.0	10.1
23439	2015.0	Lance Thomas	PF-SF	26.0	TOT	62.0	37.0	1490.0	8.0	0.456	0.050	0.224	4.8	9.9

In [134...

df[df['Pos'] == 'PG-SG']

Out[134...

	Year	Player	Pos	Age	Tm	G	GS	MP	PER	TS%	3PAr	FTr	ORB%	DRB%
20164	2010.0	Eddie House	PG-SG	31.0	TOT	68.0	0.0	1217.0	10.5	0.495	0.499	0.114	1.3	9.3
20705	2011.0	Raymond Felton	PG-SG	26.0	TOT	75.0	54.0	2737.0	16.6	0.524	0.327	0.240	2.0	9.2
22020	2013.0	Eric Maynor	PG-SG	25.0	TOT	64.0	0.0	963.0	9.3	0.472	0.342	0.221	1.1	4.4
22080	2013.0	Jeremy Pargo	PG-SG	26.0	TOT	39.0	11.0	655.0	10.4	0.478	0.296	0.248	1.3	7.5
23193	2015.0	Brandon Knight	PG-SG	23.0	TOT	63.0	61.0	2035.0	17.1	0.543	0.361	0.251	1.6	12.0
23356	2015.0	Austin Rivers	PG-SG	22.0	TOT	76.0	5.0	1563.0	10.3	0.481	0.264	0.254	2.0	8.9

In [135...

df[df['Pos'] == 'C-PF']

Out[135...

	Year	Player	Pos	Age	Tm	G	GS	MP	PER	TS%	3PAr	FTr	ORB%	DRB%	TI
20111	2010.0	Drew Gooden	C-PF	28.0	TOT	70.0	33.0	1755.0	16.9	0.547	0.013	0.359	13.7	21.9	
21802	2013.0	Ed Davis	C-PF	23.0	TOT	81.0	28.0	1631.0	17.8	0.561	0.000	0.346	10.9	22.8	

In [136...

df[df['Pos'] == 'SG-SF']

Out[136...

	Year	Player	Pos	Age	Tm	G	GS	MP	PER	TS%	3PAr	FTr	ORB%	DRB%
20459	2010.0	Henry Walker	SG-SF	22.0	TOT	35.0	13.0	768.0	14.6	0.649	0.500	0.211	2.2	11.0
20956	2011.0	Mickael Pietrus	SG-SF	28.0	TOT	57.0	4.0	1107.0	9.8	0.526	0.647	0.156	1.6	11.5
21850	2013.0	Francisco Garcia	SG-SF	31.0	TOT	58.0	20.0	1029.0	11.0	0.519	0.597	0.070	0.9	9.4
22733	2014.0	John Salmons	SG-SF	34.0	TOT	78.0	8.0	1726.0	7.8	0.462	0.395	0.126	1.3	9.6

In [137...

df[df['Pos'] == 'SG-PF']

Out[137...

	Year	Player	Pos	Age	Tm	G	GS	MP	PER	TS%	3PAr	FT%	ORB%	DRB%	T
20204	2010.0	Stephen Jackson	SG-PF	31.0	TOT	81.0	81.0	3129.0	15.6	0.518	0.277	0.306	3.0	12.2	

In [138...

df[df['Pos'] == 'PF-C']

Out[138...

	Year	Player	Pos	Age	Tm	G	GS	MP	PER	TS%	3PAr	FT%	ORB%	DRB%	T
20015	2010.0	Marcus Camby	PF-C	35.0	TOT	74.0	74.0	2314.0	17.9	0.501	0.014	0.244	12.7	31.9	
20435	2010.0	Tyrus Thomas	PF-C	23.0	TOT	54.0	3.0	1220.0	16.8	0.511	0.007	0.417	7.8	23.6	
23682	2016.0	Channing Frye	PF-C	32.0	TOT	70.0	32.0	1200.0	12.9	0.586	0.677	0.101	3.2	18.6	

In [139...

```
def change_pos(pos):
    if pos == 'PF-SF':
        return 'SF'
    if pos == 'SG-PG':
        return 'PG'
    if pos == 'PG-SG':
        return 'PG'
    if pos == 'SF-PF':
        return 'SF'
    if pos == 'C-PF':
        return 'PF'
    if pos == 'SG-SF':
        return 'SG'
    if pos == 'SF-SG':
        return 'SF'
    if pos == 'PF-C':
        return 'PF'
    if pos == 'SG-PF':
        return 'SF'
    if pos == 'C-SF':
        return 'C'
    else:
        return pos

df['Pos'] = df['Pos'].apply(lambda x : change_pos(x))
df['Pos'].value_counts()
```

Out[139...

SG	605
PG	603
SF	583
PF	573
C	481
Name: Pos, dtype: int64	


2. Merging dataset of players heights/weights

```
In [140... info = pd.read_csv("Players.csv")
info = info[['Player', 'height', 'weight']]
info.head()
```

```
Out[140...      Player  height  weight
0  Curly Armstrong   180.0    77.0
1    Cliff Barker   188.0    83.0
2   Leo Barnhorst   193.0    86.0
3    Ed Bartels   196.0    88.0
4   Ralph Beard   178.0    79.0
```

```
In [141... df = df.merge(info, on = 'Player', how = 'left')
df.head()
```

```
Out[141...      Year  Player  Pos  Age  Tm  G  GS  MP  PER  TS%  3PAr  FTr  ORB%  DRB%  TRB%
0  2010.0  Arron Afflalo  SG  24.0  DEN  82.0  75.0  2221.0  10.9  0.576  0.426  0.168  3.1  9.9  6.
1  2010.0  LaMarcus Aldridge  PF  24.0  POR  78.0  78.0  2922.0  18.2  0.535  0.014  0.260  8.1  18.6  13.
2  2010.0  Ray Allen  SG  34.0  BOS  80.0  80.0  2819.0  15.2  0.601  0.410  0.260  2.0  8.8  5.
3  2010.0  Tony Allen  SG  28.0  BOS  54.0  8.0  889.0  14.2  0.540  0.020  0.470  7.4  12.5  10.
4  2010.0  Rafer Alston  PG  33.0  TOT  52.0  38.0  1421.0  8.2  0.443  0.377  0.182  1.0  9.7  5.
```



```
In [142... df.columns
```

```
Out[142... Index(['Year', 'Player', 'Pos', 'Age', 'Tm', 'G', 'GS', 'MP', 'PER', 'TS%',
      '3PAr', 'FTr', 'ORB%', 'DRB%', 'TRB%', 'AST%', 'STL%', 'BLK%', 'TOV%',
      'USG%', 'blan1', 'OWS', 'DWS', 'WS', 'WS/48', 'blank2', 'OBPM', 'DBPM',
      'BPM', 'VORP', 'FG', 'FGA', 'FG%', '3P', '3PA', '3P%', '2P', '2PA',
      '2P%', 'eFG%', 'FT', 'FTA', 'FT%', 'ORB', 'DRB', 'TRB', 'AST', 'STL',
      'BLK', 'TOV', 'PF', 'PTS', 'MPG', 'height', 'weight'],
      dtype='object')
```

```
In [143... df.fillna(0, inplace=True)
df.isna().sum()
```

```
Out[143... Year      0
Player    0
Pos       0
Age       0
Tm        0
G         0
GS        0
MP        0
PER       0
TS%       0
3PAr      0
FTr       0
```



ORB% 0
DRB% 0
TRB% 0
AST% 0
STL% 0
BLK% 0
TOV% 0
USG% 0
blan1 0
OWS 0
DWS 0
WS 0
WS/48 0
blank2 0
OBPM 0
DBPM 0
BPM 0
VORP 0
FG 0
FGA 0
FG% 0
3P 0
3PA 0
3P% 0
2P 0
2PA 0
2P% 0
eFG% 0
FT 0
FTA 0
FT% 0
ORB 0
DRB 0
TRB 0
AST 0
STL 0
BLK 0
TOV 0
PF 0
PTS 0
MPG 0
height 0
weight 0
dtype: int64

3. Grouping by Position and finding averages by position

```
In [144... by_pos = df.groupby('Pos').mean()  
by_pos.head()
```


	Year	Age	G	GS	MP	PER	TS%	3PAr	FTr
Pos									
C	2013.569647	26.802495	61.904366	40.632017	1541.937630	16.598337	0.555112	0.037717	0.367597
PF	2013.568935	26.745201	61.783595	33.980803	1567.617801	15.379930	0.537162	0.167290	0.286969
PG	2013.570481	26.817579	60.407960	34.220564	1616.782753	14.680929	0.524381	0.330415	0.251323
SF	2013.578045	27.090909	62.049743	35.713551	1639.550600	13.015609	0.532441	0.360063	0.246926

	Year	Age	G	GS	MP	PER	TS%	3PAr	FTr
Pos									
SG	2013.528926	27.031405	61.480992	32.114050	1638.849587	13.405289	0.534098	0.377689	0.235301



```
In [145... year_pos = all_years.groupby(['Pos', 'Year'], as_index=False).mean()
year_pos = year_pos[year_pos['Year'] >= 1980]
year_pos.head()
```

```
Out[145... Pos Year Age G GS MP PER TS% 3PAr FTr
30 C 1980.0 26.830769 56.969231 41.000000 1460.076923 14.024615 0.520292 0.003277 0.360692
31 C 1981.0 26.875000 60.015625 27.333333 1458.921875 12.854687 0.511016 0.003156 0.369203
32 C 1982.0 27.106061 58.469697 29.484848 1372.075758 12.683333 0.532909 0.004212 0.409970
33 C 1983.0 27.045455 52.511364 23.795455 1170.795455 12.153409 0.502420 0.004489 0.355636
34 C 1984.0 27.478261 63.014493 30.985507 1452.072464 12.649275 0.537420 0.003232 0.376986
```




```
In [146... #only using percentages
df.columns
per = df[['Year', 'Player', 'Pos', 'Age', 'TS%', 'ORB%', 'DRB%', 'TRB%', 'AST%', 'STL%']]
```

```
In [147... non_per = df[['Year', 'Player', 'Pos', 'Age', 'PER', '3PAr', 'FTr', 'OWS', 'DWS', 'WS',
```

```
In [148... import matplotlib.pyplot as plt
import seaborn as sns
import matplotlib.gridspec as gridspec
import numpy as np
```

4. Exploratory Data Analysis

```
In [149... counts = df['Pos'].value_counts().reset_index()
plt.figure(figsize = (7,7))
sns.barplot(x = 'index', y = 'Pos', data = counts, palette = "flare")
plt.show()
```



There are significantly less centers in the dataset. This is because shooting guards and point guards make up "Guards", small forwards and power forwards make up "Forwards", but only centers are included in "Centers"

```
In [150... c = year_pos[year_pos['Pos'] == 'C']
pf = year_pos[year_pos['Pos'] == 'PF']
sf = year_pos[year_pos['Pos'] == 'SF']
sg = year_pos[year_pos['Pos'] == 'SG']
pg = year_pos[year_pos['Pos'] == 'PG']
```

```
In [151... sns.set_palette("Set1", 8, .75)

gs = gridspec.GridSpec(3,2)
plt.figure(figsize = (20,10))

ax1 = plt.subplot(gs[0, 0]) # row 0, col 0
sns.boxplot(x = 'Pos', y = '3PA', data = df)
ax1.set_title("Distribution of 3PA/Position", {'fontsize' : 18.0})
ax1.set_xlabel("Position")
ax1.set_ylabel("3PA")

ax2 = plt.subplot(gs[0, 1]) # row 0, col 1
sns.boxplot(x = 'Pos', y = '2PA', data = df)
ax2.set_title("Distribution of 2PA/Position", {'fontsize' : 18.0})
ax2.set_xlabel("Position")
ax2.set_ylabel("2PA")

plt.figure(figsize = (20,15))
ax3 = plt.subplot(gs[1,:]) # row 1, col 0
ax3.set_title("Average 3PA by Position since 1980", {'fontsize' : 18.0})
ax3.set_ylim(0,225)
```

```

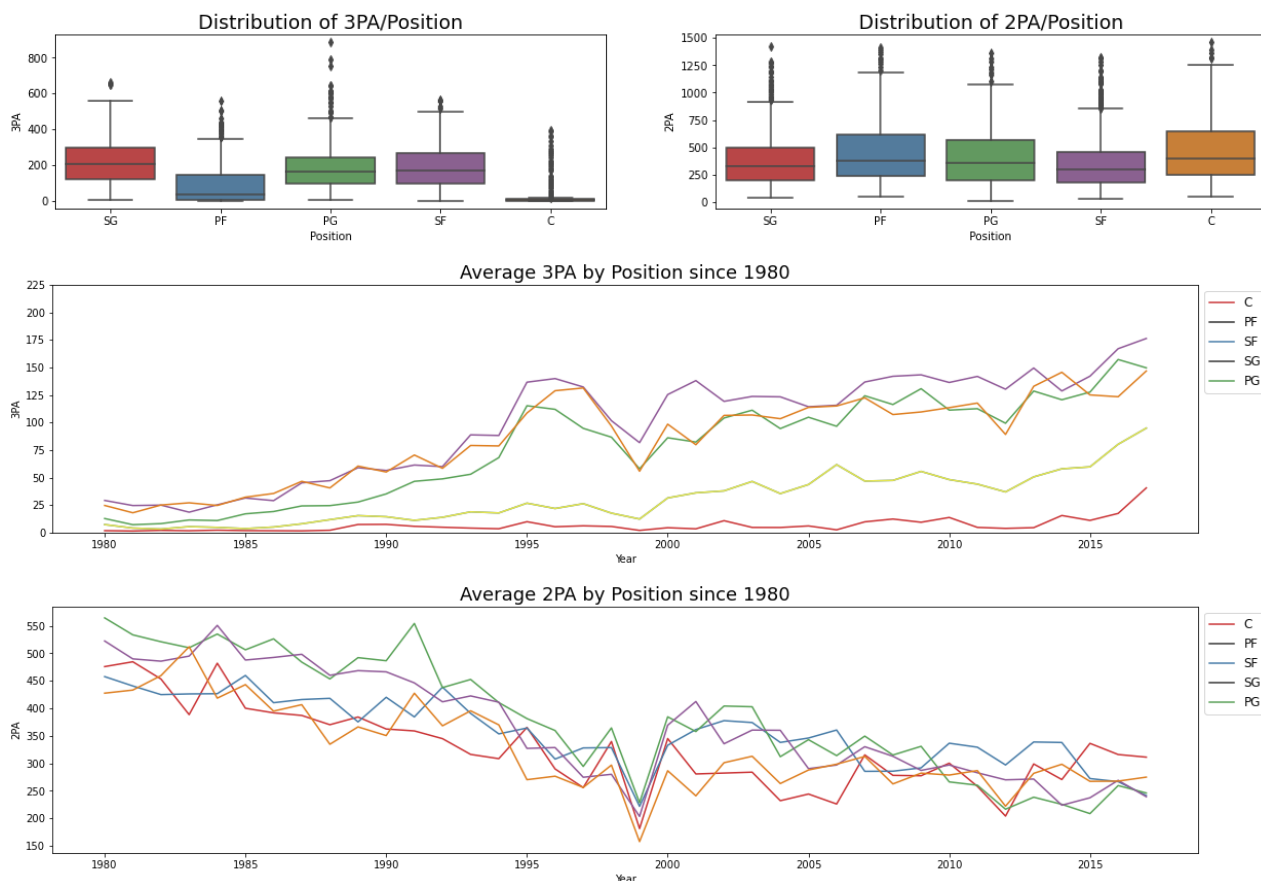
sns.lineplot(x = 'Year', y = '3PA', data = c, style = 'Pos')
sns.lineplot(x = 'Year', y = '3PA', data = pf, style = 'Pos')
sns.lineplot(x = 'Year', y = '3PA', data = sf, style = 'Pos')
sns.lineplot(x = 'Year', y = '3PA', data = sg, style = 'Pos')
sns.lineplot(x = 'Year', y = '3PA', data = pg, style = 'Pos')
sns.lineplot(x = 'Year', y = '3PA', data = pf, style = 'Pos')
plt.legend(bbox_to_anchor=(1, 1), labels=['C', 'PF', 'SF', 'SG', 'PG'], loc='upper left', pr

plt.figure(figsize = (20,15))
ax4 = plt.subplot(gs[2,:])
sns.lineplot(x = 'Year', y = '2PA', data = c, style = 'Pos', label='C')
sns.lineplot(x = 'Year', y = '2PA', data = pf, style = 'Pos', label='PF')
sns.lineplot(x = 'Year', y = '2PA', data = sf, style = 'Pos', label='SF')
sns.lineplot(x = 'Year', y = '2PA', data = sg, style = 'Pos', label='SG')
sns.lineplot(x = 'Year', y = '2PA', data = pg, style = 'Pos', label='PG')
ax4.set_title("Average 2PA by Position since 1980", {'fontsize' : 18.0})

plt.legend(bbox_to_anchor=(1, 1), labels=['C', 'PF', 'SF', 'SG', 'PG'], loc='upper left', pr

```

Out[151... <matplotlib.legend.Legend at 0x2ea2d340>




- Guards shoot more threes than forwards or centers
- Centers and power forwards shoot more two pointers than guards or small forwards
- The average 3PA by position has increased significantly since 1980 for every position
- Conversely, the average 2PA by position has decreased since 1980 for every position

```
In [152... gs = gridspec.GridSpec(1,2)

plt.figure(figsize = (20,10))
ax1 = plt.subplot(gs[0,0])
sns.boxplot(x = 'Pos', y = 'AST%', data = df)

ax2 = plt.subplot(gs[0,1])
sns.boxplot(x = 'Pos', y = 'STL%', data = df)
```

```
Out[152... <AxesSubplot:xlabel='Pos', ylabel='STL%>
```



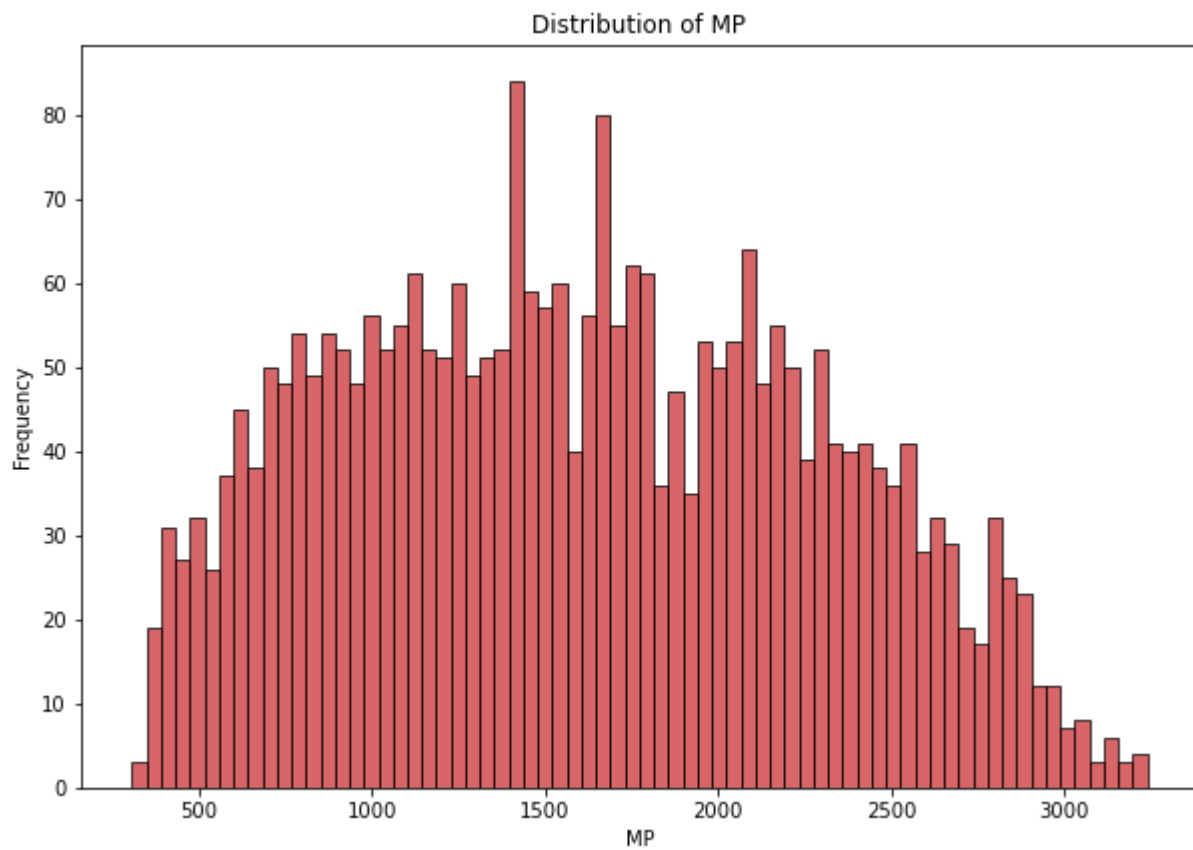

- Point guards have a much higher AST% and STL% than any other position
- Guards in general have a higher AST% and STL% than any position because they spend the most time with the ball in their hands, and defending other players that have the ball in their hands

```
In [153... gs = gridspec.GridSpec(nrows = 2,ncols = 1)

plt.figure(figsize = (10,15))
ax1 = plt.subplot(gs[0,0])
sns.histplot(x='FGA',bins=80, data = df, kde=False, ax = ax1)
ax1.set_ylabel("Frequency")
ax1.set_title("Distribution of FGA")

plt.figure(figsize = (10,15))
ax2 = plt.subplot(gs[1,0])
sns.histplot(x='MP', bins = 70, data = df, kde=False, ax = ax2)
ax2.set_ylabel("Frequency")
ax2.set_title("Distribution of MP")
```

```
Out[153... Text(0.5, 1.0, 'Distribution of MP')
```




- FGA (field goal attempts) follows a unimodal distribution, with most players shooting around 400-500 times per season
- MP (minutes played) follows a bimodal distribution, but it is almost normally distributed

```
In [154... gs = gridspec.GridSpec(2,2)

plt.figure(figsize = (15,15))
ax1 = plt.subplot(gs[0,0])
sns.boxplot(x='Pos', y='height', data = df)


ax2 = plt.subplot(gs[0,1])
sns.boxplot(x = 'Pos', y='weight', data = df)

plt.show()
```



```
In [155... gs = gridspec.GridSpec(1,2)

plt.figure(figsize = (12,12))
ax1 = plt.subplot(gs[0,0])
ax1 = sns.scatterplot(x='Pos', y='height', data = df)
```



- As expected, centers and forwards have higher heights and weights than guards, but there are some outliers.
- The outliers are from the many players that can play the position of guard, but are over 6'6", like Ben Simmons, or LeBron James.

In [156...

```
def temp(pos):
    if (pos == "PG") | (pos == 'SG'):
        return 'G'
    if (pos == 'SF') | (pos == 'PF'):
        return 'F'
    else:
        return 'C'

df['Pos'] = df['Pos'].apply(lambda x : temp(x))
```



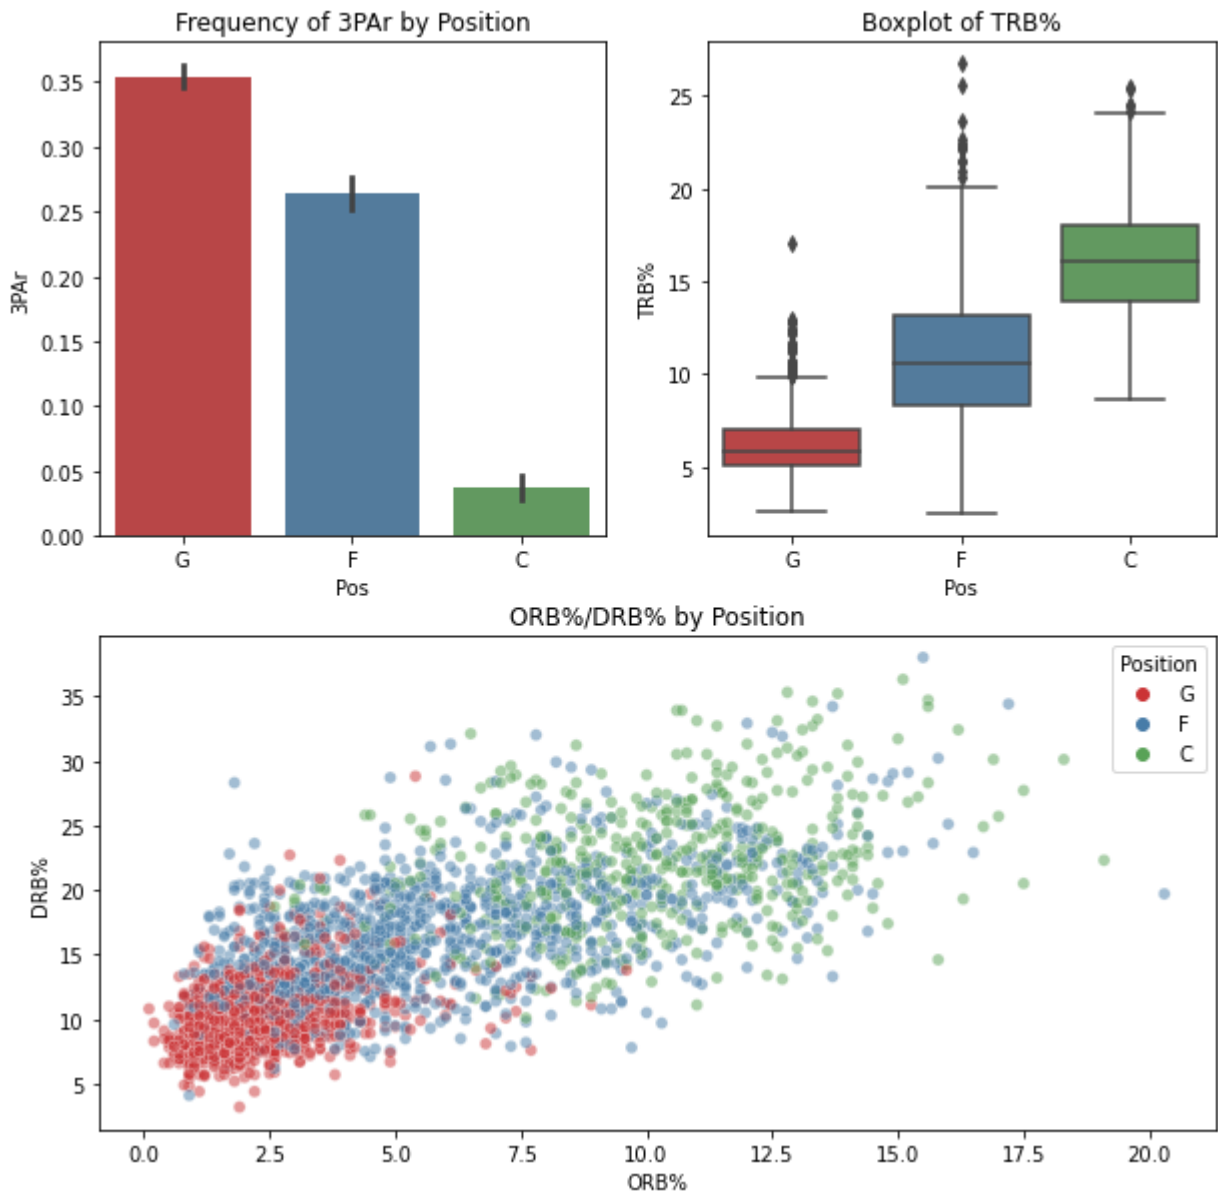
```

gs = gridspec.GridSpec(2,2)
plt.figure(figsize = (10,10))
ax1 = plt.subplot(gs[0,0])
ax1.set_title("Frequency of 3PAr by Position")
ax1 = sns.barplot(x='Pos', y= '3PAr', data=df)

ax2 = plt.subplot(gs[0,1])
ax2 = sns.boxplot(x = 'Pos', y = 'TRB%', data = df)
ax2.set_title("Boxplot of TRB%")

ax3 = plt.subplot(gs[1,:])
ax3 = sns.scatterplot(x = 'ORB%', y='DRB%', data = df, hue = 'Pos', alpha = .5)
ax3.set_title("ORB%/DRB% by Position")
plt.legend(title = "Position")
plt.show()

```



- Guards have a high relative 3PAr (3 point per field goal attempt) value
- Centers do not shoot many three pointers relative to total field goal attempts
- Centers and Forwards have a much higher TRB% than guards

- There are some outliers for both Guards and Forwards for TRB%. Again, this is because there are some guards that are taller than average, and can rebound more effectively because of their height

5. Preprocessing for Machine Learning

```
In [157... def to_encoded(pos):
    if pos == 'G':
        return 0
    if pos == 'F':
        return 1
    if pos == 'C':
        return 2

df['y'] = df['Pos'].apply(lambda x : to_encoded(x))
df['y'].value_counts()
```

```
Out[157... 0    1208
1    1156
2     481
Name: y, dtype: int64
```

Encoding Position to integers


```
In [158... #Kbest features
from sklearn.feature_selection import SelectKBest
from sklearn.feature_selection import f_classif
temp_x = df.drop(columns = ['Year', 'Player', 'Pos', 'Tm', 'G', 'GS', 'MP', 'blan1', 'bl
temp_y = df['y']

kbest = SelectKBest(score_func = f_classif, k = 15)
fit = kbest.fit(temp_x, temp_y)
dfscores = pd.DataFrame(fit.scores_)
cols = pd.DataFrame(temp_x.columns)

scores = pd.concat([cols, dfscores], axis=1)
scores.columns = ['Specs', 'Score']

scores = scores.nlargest(15, 'Score')

plt.figure(figsize = (12,12))
sns.barplot(x = 'Specs', y = 'Score', data = scores)
plt.xticks(rotation = 45)
plt.show()
```



- The most important features make sense; players with a higher rebounding percentage are more likely to be a center.
- Similarly, height and weight makes sense because taller/heavier players are more likely to be a center or forward than a guard
- AST%, 3P%, and 3PA are important because it differentiates guards from centers or forwards

```
In [159... from sklearn.metrics import confusion_matrix
def conf_mat(y_true, y_pred):
    labels = {"G" : 0,
              "F" : 1,
              "C" : 2}

    mat = confusion_matrix(y_true, y_pred)
    plot = sns.heatmap(mat, annot=True, fmt = "d", linewidths = 1, cmap = "Blues", xtic
    return plot
```

```
In [160... from sklearn.model_selection import train_test_split

X = df[['TRB%', 'ORB%', 'height', 'weight', 'FG%', 'AST%', 'BLK%', 'ORB', 'BLK', 'DBPM']
y = df['y']

X_train, X_test, y_train, y_test = train_test_split(X,y)
X_train.head()
X_train.columns
```

```
Out[160... Index(['TRB%', 'ORB%', 'height', 'weight', 'FG%', 'AST%', 'BLK%', 'ORB', 'BLK',
      'DBPM', 'AST', 'TRB', '3PA'],
      dtype='object')
```

Decision Tree Classifier with PCA

```
In [161... from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import GridSearchCV, StratifiedKFold
from sklearn.decomposition import PCA
from sklearn.cluster import KMeans
from sklearn.tree import DecisionTreeClassifier

pipe = Pipeline([
    ('std', StandardScaler()),
    ('PCA', PCA()),
    ('CLF', DecisionTreeClassifier(class_weight = 'balanced'))
])
print("PCA Parameters:", pipe['PCA'].get_params())
print()
print("KMeans paramters:", pipe['CLF'].get_params())
```

```
PCA Parameters: {'copy': True, 'iterated_power': 'auto', 'n_components': None, 'random_s
tate': None, 'svd_solver': 'auto', 'tol': 0.0, 'whiten': False}
```

```
KMeans paramters: {'ccp_alpha': 0.0, 'class_weight': 'balanced', 'criterion': 'gini', 'm
ax_depth': None, 'max_features': None, 'max_leaf_nodes': None, 'min_impurity_decrease':
0.0, 'min_impurity_split': None, 'min_samples_leaf': 1, 'min_samples_split': 2, 'min_wei
ght_fraction_leaf': 0.0, 'presort': 'deprecated', 'random_state': None, 'splitter': 'bes
t'}
```

```
In [162... cv = StratifiedKFold(shuffle=True)

param_grid = {"PCA__n_components" : [5, 7, 10, 12, 15] ,
              "PCA__whiten" : [True, False] ,
              "CLF__max_depth" : [3, 5, 10, 15] ,
              "CLF__criterion" : ['gini', 'entropy'] ,
              "CLF__max_features" : [5, 7, 9, 12, 15]}

grid_search = GridSearchCV(pipe,
                           param_grid,
                           verbose = 0,
                           scoring = 'accuracy',
                           cv = cv,
                           n_jobs = -1)

grid_search.fit(X_train, y_train)
base_score = grid_search.score(X_test, y_test)
print('Baseline Score:', base_score)
```

```
Baseline Score: 0.8089887640449438
```

```
In [163... grid_search.best_estimator_
```

```
Out[163... Pipeline(steps=[('std', StandardScaler()),
                        ('PCA', PCA(n_components=12, whiten=True)),
                        ('CLF',
                         DecisionTreeClassifier(class_weight='balanced', max_depth=5,
                                                max_features=12))])
```


```
In [164... gs = gridspec.GridSpec(3,2)
data = grid_search.cv_results_

plt.subplots(figsize = (15,15))
ax1 = plt.subplot(gs[0,0])
ax1 = sns.scatterplot(x = 'param_CLF__criterion', y = 'mean_test_score', data = data)
ax1.set_ylabel("Mean Score")
ax1.set_title("Gini/Entropy")
ax1.set_xlabel('')

ax2 = plt.subplot(gs[0,1], sharey=ax1)
ax2 = sns.scatterplot(x = "param_CLF__max_features", y = "mean_test_score", data = data)
ax2.set_ylabel('Mean Score')
ax2.set_title("Max Features")
ax2.set_xlabel('')

ax3 = plt.subplot(gs[1,:])
ax3 = sns.scatterplot(x = "param_CLF__max_depth", y = "mean_test_score", data = data)
ax3.set_xticks([3,5,10,15])
ax3.set_ylabel("Mean Score")
ax3.set_xlabel('Max Depth')


plt.show()
```



- Using gridsearch, we can see that the highest scoring trials use gini, max features

of 12, and max depth of 5

```
In [165... plt.figure(figsize = (10,10))  
  
ax4 = conf_mat(y_test, grid_search.predict(X_test))  
ax4.set_title("Confusion Matrix")  
  
plt.tight_layout()
```



The confusion matrix shows the algorithm had trouble discerning centers from forwards, and guards from forwards. Since the algorithms are using height and weight as a feature, it is most likely getting confused when it sees guards that are taller than average, and incorrectly classifies them as a forward.

Random Forest Classifier

```
In [167... from sklearn.ensemble import RandomForestClassifier

pipe2 = Pipeline([
    ('std', StandardScaler()),
    ("rf" , RandomForestClassifier(n_jobs = -1, class_weight = 'balanced'))
])

param_grid2 = {'rf__n_estimators' : [200, 250],
               'rf__max_depth' : [10, 12, 15, 17 ,20],
               'rf__min_samples_split' : [5, 7, 10]}

rf_gs = GridSearchCV(pipe2,
                     param_grid2,
                     cv = cv,
                     scoring = 'accuracy'
)

rf_gs.fit(X_train, y_train)
rf_score = rf_gs.score(X_test, y_test)
print("Baseline Score:", rf_score)
```

Baseline Score: 0.8946629213483146

```
In [168... rf_gs.best_estimator_
```

```
Out[168... Pipeline(steps=[('std', StandardScaler()),
                        ('rf',
                         RandomForestClassifier(class_weight='balanced', max_depth=20,
                                                min_samples_split=5, n_estimators=200,
                                                n_jobs=-1))])
```

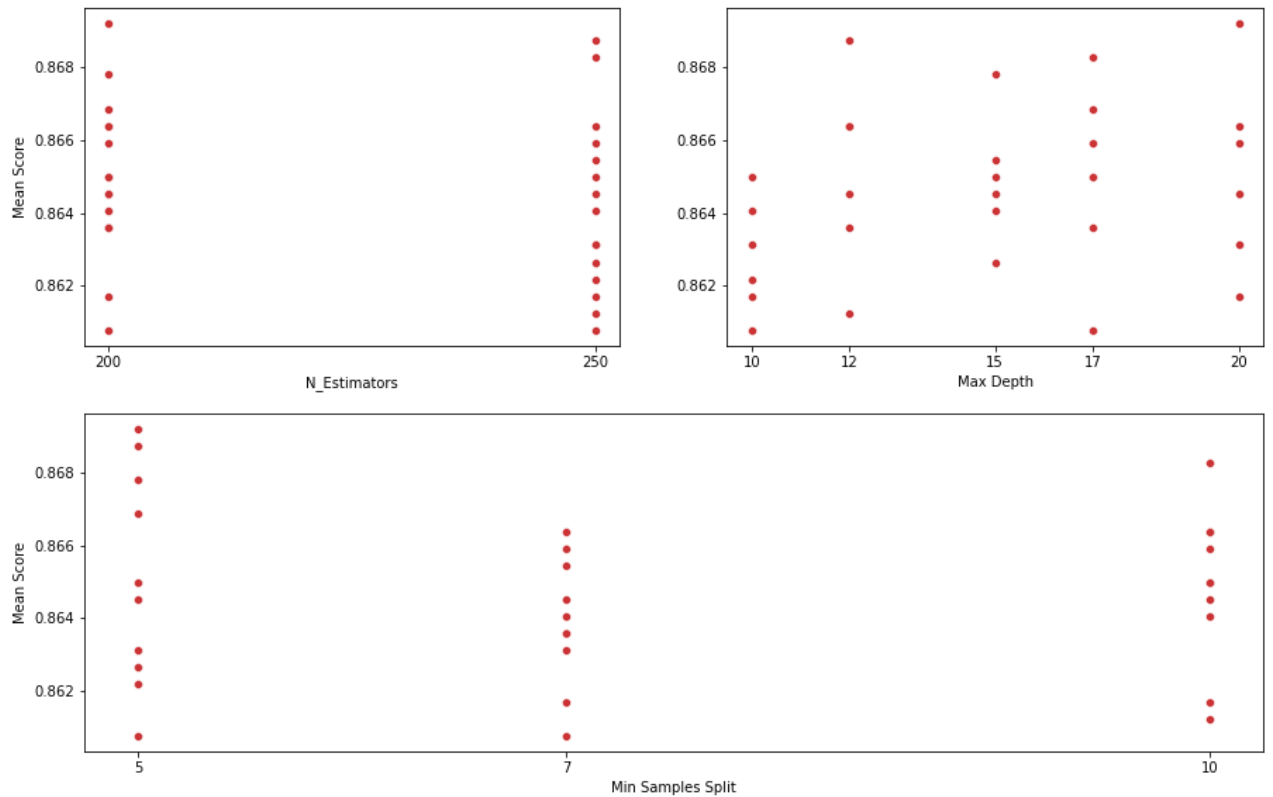
```
In [169... gs = gridspec.GridSpec(3,2)
data = rf_gs.cv_results_

plt.figure(figsize = (15,15))
ax1 = plt.subplot(gs[0,0])
sns.scatterplot(x = 'param_rf__n_estimators', y = 'mean_test_score', data = data)
ax1.set_xticks([200, 250])
ax1.set_ylabel("Mean Score")
ax1.set_xlabel("N_Estimators")

ax2 = plt.subplot(gs[0,1])
ax2 = sns.scatterplot(x = "param_rf__max_depth", y = "mean_test_score", data = data)
ax2.set_xticks([10,12,15,17, 20])
ax2.set_ylabel('')
ax2.set_xlabel('Max Depth')

ax3 = plt.subplot(gs[1,:])
ax3 = sns.scatterplot(x = "param_rf__min_samples_split", y = "mean_test_score", data = data)
ax3.set_xticks([5,7,10])
ax3.set_ylabel("Mean Score")
ax3.set_xlabel("Min Samples Split")


plt.show()
```



- Using gridsearch, we can see that the highest scoring trials use n_estimators of 200, max depth of 20, and min samples split of 5.

```
In [170... plt.figure(figsize = (10,10))

ax4 = conf_mat(y_test, rf_gs.predict(X_test))
ax4.set_title("Confusion Matrix")
plt.show()
```

Random Forest seems to be the most effective at differentiating between centers and forwards.

K-Nearest Neighbors

```
In [171... from sklearn.neighbors import KNeighborsClassifier

pipe3 = Pipeline([
    ('std', StandardScaler()),
    ('knn', KNeighborsClassifier())
])

param_grid3 = {"knn__n_neighbors" : np.arange(3, 26, 3),
               "knn__p" : [1, 2]
}

knn_gs = GridSearchCV(pipe3,
                      param_grid3,
                      cv=cv,
                      scoring='accuracy')
```

```
knn_gs.fit(X_train, y_train)
knn_score = knn_gs.score(X_test, y_test)
print("Baseline Score:", knn_score)
```

Baseline Score: 0.8707865168539326

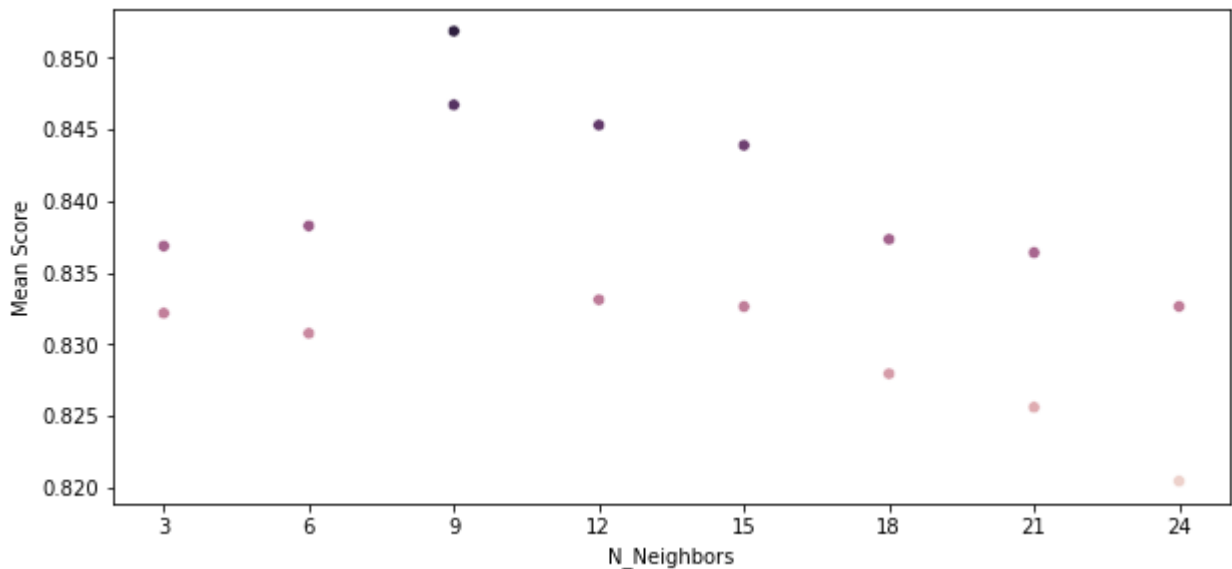
In [172... knn_gs.best_estimator_


Out[172... Pipeline(steps=[('std', StandardScaler()),
('knn', KNeighborsClassifier(n_neighbors=9, p=1))])

```
In [173... gs = gridspec.GridSpec(2,1)
data = knn_gs.cv_results_

plt.figure(figsize = (10,10))
ax1 = plt.subplot(gs[0,0])
sns.scatterplot(x = 'param_knn__n_neighbors', y = 'mean_test_score', data = data, hue="")
ax1.set_xticks(np.arange(3, 26, 3))
ax1.set_ylabel('Mean Score')
ax1.set_xlabel('N_Neighbors')
plt.legend([],[], frameon=False)

plt.figure(figsize = (15,20))
ax2 = plt.subplot(gs[1,0])
ax2.set_title("Confusion Matrix")
ax2 = conf_mat(y_test, knn_gs.predict(X_test))
```





```
In [174... from sklearn.svm import SVC
from sklearn.preprocessing import MinMaxScaler
pipe4 = Pipeline([
    ('min_max_scale', MinMaxScaler()),
    ('svm', SVC(class_weight = 'balanced', gamma = 'scale', probability=True))
])

param_grid4 = {
    'svm_C': np.arange(3, 15, 1),
    'svm_kernel' : ['rbf', 'poly']
}

svc_gs = GridSearchCV(pipe4,
                      param_grid4,
                      scoring='accuracy',
                      cv=cv,
                      n_jobs=-1)

svc_gs.fit(X_train, y_train)
svc_score = svc_gs.score(X_test, y_test)
print("Baseline Score:", svc_score)
```

Baseline Score: 0.8595505617977528

```
In [175... svc_gs.best_estimator_
```


```
Out[175... Pipeline(steps=[('min_max_scale', MinMaxScaler()),
                      ('svm', SVC(C=3, class_weight='balanced', probability=True))])
```

```
In [176... gs = gridspec.GridSpec(3,2)
data = svc_gs.cv_results_
```

```
plt.figure(figsize = (15,15))
ax1 = plt.subplot(gs[0,0])
sns.scatterplot(x = 'param_svm_C', y = 'mean_test_score', data = data)
ax1.set_xticks(np.arange(3,15,1))

ax2 = plt.subplot(gs[0,1])
ax2 = sns.scatterplot(x = "param_svm_kernel", y = "mean_test_score", data = data)

plt.figure(figsize = (20,25))
ax3 = plt.subplot(gs[1,:])
ax3 = conf_mat(y_test, svc_gs.predict(X_test))
```



In [177... `## model comparison ##`

```
In [178... models = [('Decision Tree', grid_search),
              ("Random Forest", rf_gs),
              ('KNN', knn_gs),
              ('SVM', svc_gs)]
def get_scores(models, y_test):
    temp = []
    for name, gs in models:
        fit = gs.cv_results_['mean_fit_time']
        fit = fit[~np.isnan(fit)]

        test_score = gs.cv_results_['mean_test_score']
        test_score = test_score[~np.isnan(test_score)]


        temp.append((name, gs.best_score_, test_score.mean(), fit.mean()))
```

```
return temp
```

```
model_scores = get_scores(models, y_test)  
scores = pd.DataFrame(model_scores, columns = ['Model', 'Best Score', 'Mean Test Score'])
```


In [187...

```
fig = plt.figure(figsize = (10,10))  
sns.barplot(x = 'Model', y = 'Best Score', data = scores)  
plt.ylabel('Score')  
plt.title('Best Score for each Model')  
plt.show()
```



In [183...

```
fig = plt.figure(figsize = (10,10))  
sns.barplot(x = 'Model', y = 'Mean Test Score', data = scores)  
plt.ylabel('Score')  
plt.title('Mean Test Score')  
plt.show()
```



```
In [188... fig = plt.figure(figsize = (10,10))
sns.barplot(x = 'Model', y = 'Mean Fit Time', data = scores)
plt.ylabel('Score')
plt.title('Best Score for each Model')
plt.show()
```

